

# Capstone Project

October 6, 2020

## 1 Capstone Project

### 1.1 Image classifier for the SVHN dataset

#### 1.1.1 Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

#### 1.1.2 How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

#### 1.1.3 Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
In [1]: import math
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import gridspec
from scipy.io import loadmat

# Tensorflow and Keras imports
import tensorflow as tf
from tensorflow.keras.models import load_model
```

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, LeakyReLU
from tensorflow.keras.layers import BatchNormalization, Dropout
from tensorflow.keras.layers import Conv2D, MaxPooling2D

from tensorflow.keras import regularizers
from tensorflow.keras.callbacks import LearningRateScheduler
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.callbacks import ModelCheckpoint

from sklearn.preprocessing import MultiLabelBinarizer
from sklearn.metrics import confusion_matrix

```



For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. “Reading Digits in Natural Images with Unsupervised Feature Learning”. NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

In [2]: *# Run this cell to load the dataset*

```

train = loadmat('data/train_32x32.mat')
test = loadmat('data/test_32x32.mat')

```

Both train and test are dictionaries with keys X and y for the input images and labels respectively.

## 1.2 1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.


```
In [3]: X_train, X_test = train['X'], test['X']  
        y_train, y_test = train['y'], test['y']
```

```
        new_order = [3,0,1,2]  
        X_train = X_train.transpose(new_order)  
        X_test = X_test.transpose(new_order)
```

```
In [4]: def rectify_labels(old_labels):  
        corrected_labels = np.where(old_labels == 10, 0, old_labels) # replace 10 by 0  
        return corrected_labels
```

```
In [5]: y_train = rectify_labels(y_train)  
        y_test = rectify_labels(y_test)
```

```
In [6]: plt.imshow(X_train[35,:,:,:])  
        img_id = 'This is the number ' + str(y_train[35][0])  
        plt.title(img_id)  
        plt.show()
```



```

In [7]: def show_images(X_data, y_data, num_figs):
    num_examples = X_data.shape[0]
    sample = np.random.choice(range(num_examples), size=num_figs)

    n_rows = int(math.ceil(num_figs / 10))
    gs = gridspec.GridSpec(n_rows, 10)
    fig = plt.figure(figsize=(18,4))

    for i,j in enumerate(sample):
        if X_data.shape[-1] == 3:
            image = X_data[j,:,:, :]
        elif X_data.shape[-1] == 1:
            image = X_data[j,:,:,0]
        else:
            raise ValueError('Wrong dimension of the input data...')
        img_label = 'Number ' + str(y_data[j][0])
        ax = fig.add_subplot(gs[i])
        ax.set_axis_off()
        ax.imshow(image)
        ax.set_title(img_label)

    plt.show()

```

```

In [8]: show_images(X_train, y_train, 15)

```




```

In [9]: def average_colors(data):
    data_grayscale = np.average(data, axis = -1)[...,np.newaxis]
    return data_grayscale

In [10]: X_train = average_colors(X_train)
    X_test = average_colors(X_test)

In [11]: show_images(X_train, y_train, 10)

```



```
In [12]: X_train = X_train / 255.0
        X_test = X_test / 255.0
```

### 1.3 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the `summary()` method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a `ModelCheckpoint` callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [13]: def leaky_relu(x):
        return tf.nn.leaky_relu(x, alpha=0.15)

        def get_model(neurons_per_layers, wd, train_data):
            model = Sequential()
            for i, n_neurons in enumerate(neurons_per_layers):
                if i == 0:
                    model.add(Flatten(input_shape=train_data.shape[1:]))
                    model.add(Dense(n_neurons, kernel_regularizer=regularizers.l2(wd),
                                    activation=leaky_relu))
                else:
                    model.add(Dense(n_neurons, kernel_regularizer=regularizers.l2(wd),
                                    activation=leaky_relu))

            model.add(Dense(10, activation='softmax', name="out_layer"))
            return model

In [14]: nn_arch = [100, 50]
        weight_decay = 1e-6
        model = get_model(nn_arch, weight_decay, X_train)
        model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 100)	102500
dense_1 (Dense)	(None, 50)	5050
out_layer (Dense)	(None, 10)	510
Total params: 108,060		
Trainable params: 108,060		
Non-trainable params: 0		

```
In [15]: opt = tf.keras.optimizers.Adam(learning_rate=0.01)
        loss = tf.keras.losses.SparseCategoricalCrossentropy()

        model.compile(optimizer = opt, loss = loss, metrics = ['accuracy'])

In [16]: def lr_decay(epoch, lr):
        if epoch <= 3:
            return lr
        else:
            lr = lr * np.exp(-1.5/epoch)
            return lr

        # learning schedule callback
        lrate = LearningRateScheduler(lr_decay, verbose=1)
        earlystop = EarlyStopping(patience=3)
        mlp_checkpoint_best = 'nn_checkpoints_best/checkpoint'
        checkpoint_best = ModelCheckpoint(filepath=mlp_checkpoint_best, save_freq='epoch',
                                          save_weights_only=False, monitor='val_accuracy',
                                          save_best_only=True, verbose=1)

        my_callbacks = [lrate, earlystop, checkpoint_best]

In [17]: history = model.fit(X_train, y_train, epochs = 30, validation_split=0.15,
                            batch_size = 100, callbacks=my_callbacks, verbose = 1)

Train on 62268 samples, validate on 10989 samples

Epoch 00001: LearningRateScheduler reducing learning rate to 0.009999999776482582.
Epoch 1/30
61900/62268 [=====>.] - ETA: 0s - loss: 2.3114 - accuracy: 0.1861
Epoch 00001: val_accuracy improved from -inf to 0.21704, saving model to nn_checkpoints_best/cl
```

```

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/resonance_ops.py:110:
Instructions for updating:
If using Keras pass *_constraint arguments to layers.
INFO:tensorflow:Assets written to: nn_checkpoints_best/checkpoint/assets
62268/62268 [=====] - 19s 312us/sample - loss: 2.3108 - accuracy: 0.1875

Epoch 00002: LearningRateScheduler reducing learning rate to 0.009999999776482582.
Epoch 2/30
62200/62268 [=====>.] - ETA: 0s - loss: 1.9402 - accuracy: 0.3119 ETA: 0s
Epoch 00002: val_accuracy improved from 0.21704 to 0.43116, saving model to nn_checkpoints_best/checkpoint/assets
INFO:tensorflow:Assets written to: nn_checkpoints_best/checkpoint/assets
62268/62268 [=====] - 17s 270us/sample - loss: 1.9400 - accuracy: 0.3119

Epoch 00003: LearningRateScheduler reducing learning rate to 0.009999999776482582.
Epoch 3/30
62000/62268 [=====>.] - ETA: 0s - loss: 1.4946 - accuracy: 0.5008
Epoch 00003: val_accuracy improved from 0.43116 to 0.56183, saving model to nn_checkpoints_best/checkpoint/assets
INFO:tensorflow:Assets written to: nn_checkpoints_best/checkpoint/assets
62268/62268 [=====] - 17s 275us/sample - loss: 1.4935 - accuracy: 0.5008

Epoch 00004: LearningRateScheduler reducing learning rate to 0.009999999776482582.
Epoch 4/30
62000/62268 [=====>.] - ETA: 0s - loss: 1.3262 - accuracy: 0.5732
Epoch 00004: val_accuracy improved from 0.56183 to 0.59077, saving model to nn_checkpoints_best/checkpoint/assets
INFO:tensorflow:Assets written to: nn_checkpoints_best/checkpoint/assets
62268/62268 [=====] - 17s 279us/sample - loss: 1.3263 - accuracy: 0.5732

Epoch 00005: LearningRateScheduler reducing learning rate to 0.006872892634288598.
Epoch 5/30
62200/62268 [=====>.] - ETA: 0s - loss: 1.1849 - accuracy: 0.6300
Epoch 00005: val_accuracy improved from 0.59077 to 0.63008, saving model to nn_checkpoints_best/checkpoint/assets
INFO:tensorflow:Assets written to: nn_checkpoints_best/checkpoint/assets
62268/62268 [=====] - 17s 276us/sample - loss: 1.1848 - accuracy: 0.6300

Epoch 00006: LearningRateScheduler reducing learning rate to 0.005091564248681657.
Epoch 6/30
62000/62268 [=====>.] - ETA: 0s - loss: 1.0912 - accuracy: 0.6665
Epoch 00006: val_accuracy improved from 0.63008 to 0.65283, saving model to nn_checkpoints_best/checkpoint/assets
INFO:tensorflow:Assets written to: nn_checkpoints_best/checkpoint/assets
62268/62268 [=====] - 17s 279us/sample - loss: 1.0907 - accuracy: 0.6528

Epoch 00007: LearningRateScheduler reducing learning rate to 0.003965314235977213.
Epoch 7/30
62100/62268 [=====>.] - ETA: 0s - loss: 1.0200 - accuracy: 0.6938
Epoch 00007: val_accuracy improved from 0.65283 to 0.68405, saving model to nn_checkpoints_best/checkpoint/assets
INFO:tensorflow:Assets written to: nn_checkpoints_best/checkpoint/assets
62268/62268 [=====] - 17s 274us/sample - loss: 1.0200 - accuracy: 0.6840

```

Epoch 00008: LearningRateScheduler reducing learning rate to 0.0032004753115511507.  
Epoch 8/30  
61900/62268 [=====>.] - ETA: 0s - loss: 0.9602 - accuracy: 0.7107  
Epoch 00008: val\_accuracy improved from 0.68405 to 0.70962, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 277us/sample - loss: 0.9599 - accuracy: 0.7096

Epoch 00009: LearningRateScheduler reducing learning rate to 0.002653287264908044.  
Epoch 9/30  
62100/62268 [=====>.] - ETA: 0s - loss: 0.9099 - accuracy: 0.7285  
Epoch 00009: val\_accuracy improved from 0.70962 to 0.71763, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 273us/sample - loss: 0.9098 - accuracy: 0.7176

Epoch 00010: LearningRateScheduler reducing learning rate to 0.0022459591747814355.  
Epoch 10/30  
62100/62268 [=====>.] - ETA: 0s - loss: 0.8721 - accuracy: 0.7419  
Epoch 00010: val\_accuracy improved from 0.71763 to 0.73346, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 18s 285us/sample - loss: 0.8721 - accuracy: 0.7334

Epoch 00011: LearningRateScheduler reducing learning rate to 0.0019331149356105692.  
Epoch 11/30  
61800/62268 [=====>.] - ETA: 0s - loss: 0.8398 - accuracy: 0.7487  
Epoch 00011: val\_accuracy improved from 0.73346 to 0.73756, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 274us/sample - loss: 0.8399 - accuracy: 0.7375

Epoch 00012: LearningRateScheduler reducing learning rate to 0.0016866916899998582.  
Epoch 12/30  
62000/62268 [=====>.] - ETA: 0s - loss: 0.8054 - accuracy: 0.7604  
Epoch 00012: val\_accuracy improved from 0.73756 to 0.75576, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 278us/sample - loss: 0.8054 - accuracy: 0.7557

Epoch 00013: LearningRateScheduler reducing learning rate to 0.0014885001618388336.  
Epoch 13/30  
62200/62268 [=====>.] - ETA: 0s - loss: 0.7757 - accuracy: 0.7697  
Epoch 00013: val\_accuracy did not improve from 0.75576  
62268/62268 [=====] - 16s 250us/sample - loss: 0.7757 - accuracy: 0.7557

Epoch 00014: LearningRateScheduler reducing learning rate to 0.0013262884372625512.  
Epoch 14/30  
62100/62268 [=====>.] - ETA: 0s - loss: 0.7505 - accuracy: 0.7776 ETA: 0s  
Epoch 00014: val\_accuracy did not improve from 0.75576  
62268/62268 [=====] - 15s 241us/sample - loss: 0.7508 - accuracy: 0.7557

Epoch 00015: LearningRateScheduler reducing learning rate to 0.001191533987088959.



Epoch 15/30  
62000/62268 [=====>.] - ETA: 0s - loss: 0.7246 - accuracy: 0.7853  
Epoch 00015: val\_accuracy improved from 0.75576 to 0.77368, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 16s 264us/sample - loss: 0.7251 - accuracy: 0.7853

Epoch 00016: LearningRateScheduler reducing learning rate to 0.0010781445349850964.  
Epoch 16/30  
62200/62268 [=====>.] - ETA: 0s - loss: 0.7076 - accuracy: 0.7900  
Epoch 00016: val\_accuracy did not improve from 0.77368  
62268/62268 [=====] - 15s 245us/sample - loss: 0.7074 - accuracy: 0.7900

Epoch 00017: LearningRateScheduler reducing learning rate to 0.000981661766327396.  
Epoch 17/30  
61900/62268 [=====>.] - ETA: 0s - loss: 0.6906 - accuracy: 0.7955  
Epoch 00017: val\_accuracy improved from 0.77368 to 0.77641, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 268us/sample - loss: 0.6899 - accuracy: 0.7955

Epoch 00018: LearningRateScheduler reducing learning rate to 0.0008987559040903719.  
Epoch 18/30  
62100/62268 [=====>.] - ETA: 0s - loss: 0.6753 - accuracy: 0.8009 ETA: 0s  
Epoch 00018: val\_accuracy improved from 0.77641 to 0.78488, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 271us/sample - loss: 0.6750 - accuracy: 0.8009

Epoch 00019: LearningRateScheduler reducing learning rate to 0.000826895357137767.  
Epoch 19/30  
62000/62268 [=====>.] - ETA: 0s - loss: 0.6590 - accuracy: 0.8064  
Epoch 00019: val\_accuracy improved from 0.78488 to 0.78688, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 270us/sample - loss: 0.6589 - accuracy: 0.8064

Epoch 00020: LearningRateScheduler reducing learning rate to 0.0007641245358877472.  
Epoch 20/30  
61900/62268 [=====>.] - ETA: 0s - loss: 0.6489 - accuracy: 0.8088 ETA: 0s  
Epoch 00020: val\_accuracy improved from 0.78688 to 0.79134, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 269us/sample - loss: 0.6490 - accuracy: 0.8088

Epoch 00021: LearningRateScheduler reducing learning rate to 0.0007089115843053009.  
Epoch 21/30  
61900/62268 [=====>.] - ETA: 0s - loss: 0.6355 - accuracy: 0.8134  
Epoch 00021: val\_accuracy improved from 0.79134 to 0.79361, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 277us/sample - loss: 0.6352 - accuracy: 0.8134

Epoch 00022: LearningRateScheduler reducing learning rate to 0.0006600411749978975.

Epoch 22/30  
62000/62268 [=====>.] - ETA: 0s - loss: 0.6284 - accuracy: 0.8156 ETA: 0s  
Epoch 00022: val\_accuracy improved from 0.79361 to 0.79743, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 265us/sample - loss: 0.6279 - accuracy: 0.8156

Epoch 00023: LearningRateScheduler reducing learning rate to 0.0006165382703631737.  
Epoch 23/30  
62000/62268 [=====>.] - ETA: 0s - loss: 0.6166 - accuracy: 0.8190 ETA: 0s  
Epoch 00023: val\_accuracy improved from 0.79743 to 0.79934, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 270us/sample - loss: 0.6166 - accuracy: 0.8190

Epoch 00024: LearningRateScheduler reducing learning rate to 0.0005776123807018673.  
Epoch 24/30  
61900/62268 [=====>.] - ETA: 0s - loss: 0.6086 - accuracy: 0.8212  
Epoch 00024: val\_accuracy did not improve from 0.79934  
62268/62268 [=====] - 15s 247us/sample - loss: 0.6083 - accuracy: 0.8212

Epoch 00025: LearningRateScheduler reducing learning rate to 0.0005426166080451965.  
Epoch 25/30  
62200/62268 [=====>.] - ETA: 0s - loss: 0.6008 - accuracy: 0.8237  
Epoch 00025: val\_accuracy improved from 0.79934 to 0.80116, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 17s 269us/sample - loss: 0.6010 - accuracy: 0.8237

Epoch 00026: LearningRateScheduler reducing learning rate to 0.0005110170579727344.  
Epoch 26/30  
62000/62268 [=====>.] - ETA: 0s - loss: 0.5944 - accuracy: 0.8265  
Epoch 00026: val\_accuracy improved from 0.80116 to 0.80144, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 16s 261us/sample - loss: 0.5942 - accuracy: 0.8265

Epoch 00027: LearningRateScheduler reducing learning rate to 0.0004823696071240638.  
Epoch 27/30  
61900/62268 [=====>.] - ETA: 0s - loss: 0.5880 - accuracy: 0.8278  
Epoch 00027: val\_accuracy improved from 0.80144 to 0.80653, saving model to nn\_checkpoints\_best  
INFO:tensorflow:Assets written to: nn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 16s 259us/sample - loss: 0.5883 - accuracy: 0.8278

Epoch 00028: LearningRateScheduler reducing learning rate to 0.0004563020836763124.  
Epoch 28/30  
61800/62268 [=====>.] - ETA: 0s - loss: 0.5819 - accuracy: 0.8302  
Epoch 00028: val\_accuracy did not improve from 0.80653  
62268/62268 [=====] - 15s 245us/sample - loss: 0.5818 - accuracy: 0.8302

Epoch 00029: LearningRateScheduler reducing learning rate to 0.0004325005561633849.  
Epoch 29/30

```

61800/62268 [=====>.] - ETA: 0s - loss: 0.5755 - accuracy: 0.8319
Epoch 00029: val_accuracy did not improve from 0.80653
62268/62268 [=====] - 15s 249us/sample - loss: 0.5762 - accuracy: 0.8319

Epoch 00030: LearningRateScheduler reducing learning rate to 0.0004106985446929329.
Epoch 30/30
61900/62268 [=====>.] - ETA: 0s - loss: 0.5712 - accuracy: 0.8320 ETA: 0s
Epoch 00030: val_accuracy improved from 0.80653 to 0.80917, saving model to nn_checkpoints_best
INFO:tensorflow:Assets written to: nn_checkpoints_best/checkpoint/assets
62268/62268 [=====] - 17s 268us/sample - loss: 0.5710 - accuracy: 0.8320

```

```

In [18]: df_hist = pd.DataFrame(history.history)
         df_hist.head(10)

```

```

Out[18]:
```

	loss	accuracy	val_loss	val_accuracy	lr
0	2.310797	0.186131	2.172417	0.217035	0.010000
1	1.939972	0.312070	1.629429	0.431158	0.010000
2	1.493482	0.501140	1.367031	0.561835	0.010000
3	1.326325	0.573232	1.309591	0.590773	0.010000
4	1.184806	0.630019	1.205127	0.630085	0.006873
5	1.090663	0.666602	1.111646	0.652835	0.005092
6	1.019994	0.693791	1.047261	0.684048	0.003965
7	0.959919	0.710863	0.970394	0.709619	0.003200
8	0.909784	0.728448	0.926577	0.717627	0.002653
9	0.872070	0.741874	0.889884	0.733461	0.002246

```

In [19]: def plot_results(history):
         fig = plt.figure(figsize=(12, 5))

         fig.add_subplot(121)

         plt.plot(history['loss'])
         plt.plot(history['val_loss'])
         plt.title('Loss vs. epochs')
         plt.ylabel('Loss')
         plt.xlabel('Epoch')
         plt.legend(['Training', 'Validation'], loc='upper right')

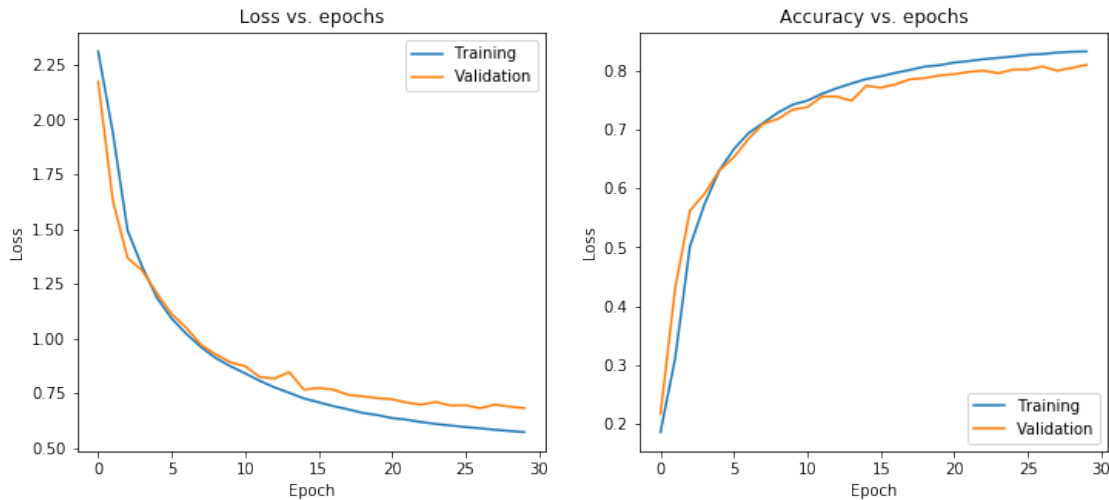
         fig.add_subplot(122)

         plt.plot(history['accuracy'])
         plt.plot(history['val_accuracy'])
         plt.title('Accuracy vs. epochs')
         plt.ylabel('Accuracy')
         plt.xlabel('Epoch')
         plt.legend(['Training', 'Validation'], loc='lower right')

         plt.show()

```

```
In [20]: plot_results(df_hist)
```



```
In [21]: def get_test_accuracy(model, X_test, y_test):  
    test_loss, test_acc = model.evaluate(x=X_test, y=y_test, verbose=0)  
    print('accuracy: {acc:0.3f}'.format(acc=test_acc))  
    print('loss: {loss:0.3f}'.format(loss=test_loss))  
  
    get_test_accuracy(model, X_test, y_test)
```

accuracy: 0.793

loss: 0.781

### 1.4 3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.)*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```

In [22]: def get_cnn_model(train_data, wd, dprate, lr):
    model = Sequential()
    # Convolutional layers
    model.add(Conv2D(filters=24, kernel_size=(3,3), padding='SAME', activation='relu',
                      input_shape = train_data.shape[1:], name='conv_1'))
    model.add(MaxPooling2D(pool_size=(2,2)))
    model.add(BatchNormalization())
    model.add(Dropout(dprate))

    model.add(Conv2D(filters=16, kernel_size=(3,3), padding='SAME',
                      activation='relu', name='conv_2'))
    model.add(MaxPooling2D(pool_size=(3,3)))
    model.add(BatchNormalization())
    model.add(Dropout(dprate*0.5))

    model.add(Flatten())
    # Fully connected layers
    model.add(Dense(50, activation=leaky_relu, kernel_regularizer=regularizers.l2(wd))

    # Output layer
    model.add(Dense(10, activation='softmax'))

    params = {'optimizer': tf.keras.optimizers.Adam(learning_rate=lr),
              'loss': tf.keras.losses.SparseCategoricalCrossentropy(),
              'metrics': ['accuracy']}

    model.compile(**params)

    return model

In [23]: weight_decay = 1e-6
    dropout_rate = 0.5
    opt_lr = 0.01
    cnn_model = get_cnn_model(X_train, weight_decay, dropout_rate, opt_lr)
    cnn_model.summary()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv_1 (Conv2D)	(None, 32, 32, 24)	240
max_pooling2d (MaxPooling2D)	(None, 16, 16, 24)	0
batch_normalization (BatchNormalizatio	(None, 16, 16, 24)	96
dropout (Dropout)	(None, 16, 16, 24)	0

```

-----
conv_2 (Conv2D)                (None, 16, 16, 16)        3472
-----
max_pooling2d_1 (MaxPooling2D) (None, 5, 5, 16)          0
-----
batch_normalization_1 (Batch Normalization) (None, 5, 5, 16)          64
-----
dropout_1 (Dropout)            (None, 5, 5, 16)          0
-----
flatten_1 (Flatten)            (None, 400)                0
-----
fc_1 (Dense)                   (None, 50)                 20050
-----
dense_2 (Dense)                (None, 10)                 510
=====
Total params: 24,432
Trainable params: 24,352
Non-trainable params: 80
-----

```

```

In [24]: cnn_checkpoint_best = 'cnn_checkpoints_best/checkpoint'
         checkpoint_best = ModelCheckpoint(filepath=cnn_checkpoint_best, save_freq='epoch',
         save_weights_only=False, monitor='val_accuracy',
         save_best_only=True, verbose=1)

```

```

my_callbacks = [lr_scheduler, early_stopping, checkpoint_best]

```

```

In [25]: history = cnn_model.fit(X_train, y_train, epochs = 20, validation_split=0.15,
                                batch_size = 64, callbacks=my_callbacks, verbose = 1)

```

Train on 62268 samples, validate on 10989 samples

Epoch 00001: LearningRateScheduler reducing learning rate to 0.009999999776482582.

Epoch 1/20

62208/62268 [=====>.] - ETA: 0s - loss: 1.0542 - accuracy: 0.6558

Epoch 00001: val\_accuracy improved from -inf to 0.73273, saving model to cnn\_checkpoints\_best/

INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets

62268/62268 [=====] - 283s 5ms/sample - loss: 1.0539 - accuracy: 0.6558

Epoch 00002: LearningRateScheduler reducing learning rate to 0.009999999776482582.

Epoch 2/20

62208/62268 [=====>.] - ETA: 0s - loss: 0.7037 - accuracy: 0.7795

Epoch 00002: val\_accuracy improved from 0.73273 to 0.83347, saving model to cnn\_checkpoints\_best/

INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets

62268/62268 [=====] - 277s 4ms/sample - loss: 0.7037 - accuracy: 0.7795

Epoch 00003: LearningRateScheduler reducing learning rate to 0.009999999776482582.

Epoch 3/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.6529 - accuracy: 0.7958  
Epoch 00003: val\_accuracy improved from 0.83347 to 0.84393, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 268s 4ms/sample - loss: 0.6527 - accuracy: 0.7958

Epoch 00004: LearningRateScheduler reducing learning rate to 0.009999999776482582.  
Epoch 4/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.6411 - accuracy: 0.7985  
Epoch 00004: val\_accuracy improved from 0.84393 to 0.84739, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 257s 4ms/sample - loss: 0.6412 - accuracy: 0.7985

Epoch 00005: LearningRateScheduler reducing learning rate to 0.006872892634288598.  
Epoch 5/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.5845 - accuracy: 0.8188  
Epoch 00005: val\_accuracy improved from 0.84739 to 0.85203, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 247s 4ms/sample - loss: 0.5845 - accuracy: 0.8188

Epoch 00006: LearningRateScheduler reducing learning rate to 0.005091564248681657.  
Epoch 6/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.5550 - accuracy: 0.8274  
Epoch 00006: val\_accuracy improved from 0.85203 to 0.87524, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 256s 4ms/sample - loss: 0.5551 - accuracy: 0.8274

Epoch 00007: LearningRateScheduler reducing learning rate to 0.003965314235977213.  
Epoch 7/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.5301 - accuracy: 0.8355  
Epoch 00007: val\_accuracy improved from 0.87524 to 0.88243, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 258s 4ms/sample - loss: 0.5303 - accuracy: 0.8355

Epoch 00008: LearningRateScheduler reducing learning rate to 0.0032004753115511507.  
Epoch 8/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.5136 - accuracy: 0.8413  
Epoch 00008: val\_accuracy did not improve from 0.88243  
62268/62268 [=====] - 258s 4ms/sample - loss: 0.5135 - accuracy: 0.8413

Epoch 00009: LearningRateScheduler reducing learning rate to 0.002653287264908044.  
Epoch 9/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.4984 - accuracy: 0.8456  
Epoch 00009: val\_accuracy improved from 0.88243 to 0.88698, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 257s 4ms/sample - loss: 0.4985 - accuracy: 0.8456

Epoch 00010: LearningRateScheduler reducing learning rate to 0.0022459591747814355.

Epoch 10/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.4951 - accuracy: 0.8482  
Epoch 00010: val\_accuracy did not improve from 0.88698  
62268/62268 [=====] - 253s 4ms/sample - loss: 0.4949 - accuracy: 0.8482

Epoch 00011: LearningRateScheduler reducing learning rate to 0.0019331149356105692.  
Epoch 11/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.4894 - accuracy: 0.8492  
Epoch 00011: val\_accuracy improved from 0.88698 to 0.88998, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 255s 4ms/sample - loss: 0.4892 - accuracy: 0.8492

Epoch 00012: LearningRateScheduler reducing learning rate to 0.0016866916899998582.  
Epoch 12/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.4778 - accuracy: 0.8523  
Epoch 00012: val\_accuracy improved from 0.88998 to 0.89053, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 259s 4ms/sample - loss: 0.4781 - accuracy: 0.8523

Epoch 00013: LearningRateScheduler reducing learning rate to 0.0014885001618388336.  
Epoch 13/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.4728 - accuracy: 0.8537  
Epoch 00013: val\_accuracy improved from 0.89053 to 0.89235, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 259s 4ms/sample - loss: 0.4727 - accuracy: 0.8537

Epoch 00014: LearningRateScheduler reducing learning rate to 0.0013262884372625512.  
Epoch 14/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.4691 - accuracy: 0.8545  
Epoch 00014: val\_accuracy did not improve from 0.89235  
62268/62268 [=====] - 255s 4ms/sample - loss: 0.4690 - accuracy: 0.8545

Epoch 00015: LearningRateScheduler reducing learning rate to 0.001191533987088959.  
Epoch 15/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.4674 - accuracy: 0.8558  
Epoch 00015: val\_accuracy improved from 0.89235 to 0.89353, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 257s 4ms/sample - loss: 0.4674 - accuracy: 0.8558

Epoch 00016: LearningRateScheduler reducing learning rate to 0.0010781445349850964.  
Epoch 16/20  
62208/62268 [=====>.] - ETA: 0s - loss: 0.4608 - accuracy: 0.8586  
Epoch 00016: val\_accuracy improved from 0.89353 to 0.89489, saving model to cnn\_checkpoints\_best  
INFO:tensorflow:Assets written to: cnn\_checkpoints\_best/checkpoint/assets  
62268/62268 [=====] - 264s 4ms/sample - loss: 0.4606 - accuracy: 0.8586

Epoch 00017: LearningRateScheduler reducing learning rate to 0.000981661766327396.  
Epoch 17/20



```

62208/62268 [=====>.] - ETA: 0s - loss: 0.4606 - accuracy: 0.8574
Epoch 00017: val_accuracy did not improve from 0.89489
62268/62268 [=====] - 258s 4ms/sample - loss: 0.4606 - accuracy: 0.8574

Epoch 00018: LearningRateScheduler reducing learning rate to 0.0008987559040903719.
Epoch 18/20
62208/62268 [=====>.] - ETA: 0s - loss: 0.4613 - accuracy: 0.8579 ETA: 0s
Epoch 00018: val_accuracy did not improve from 0.89489
62268/62268 [=====] - 254s 4ms/sample - loss: 0.4614 - accuracy: 0.8579

Epoch 00019: LearningRateScheduler reducing learning rate to 0.000826895357137767.
Epoch 19/20
62208/62268 [=====>.] - ETA: 0s - loss: 0.4553 - accuracy: 0.8607
Epoch 00019: val_accuracy did not improve from 0.89489
62268/62268 [=====] - 258s 4ms/sample - loss: 0.4551 - accuracy: 0.8607

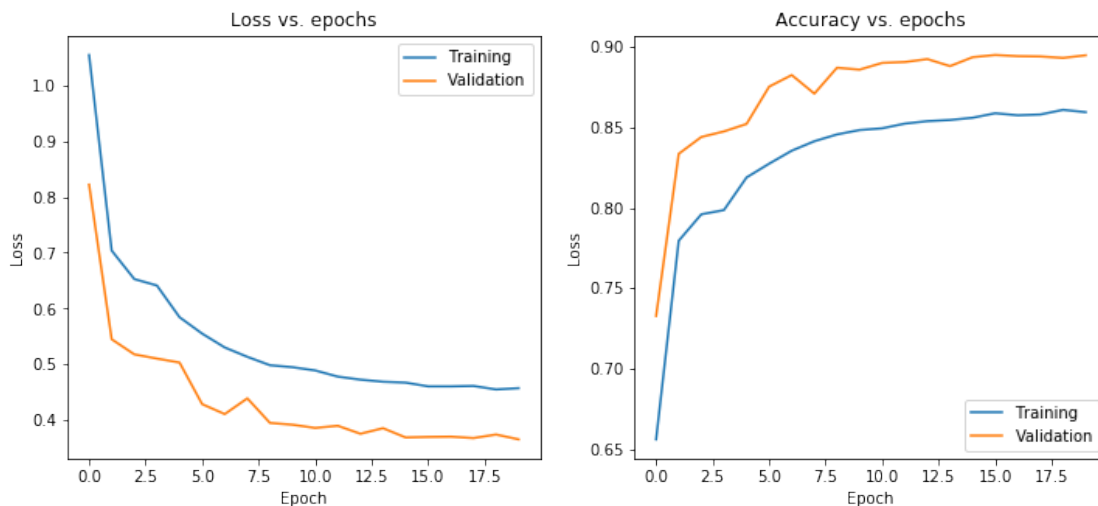
Epoch 00020: LearningRateScheduler reducing learning rate to 0.0007641245358877472.
Epoch 20/20
62208/62268 [=====>.] - ETA: 0s - loss: 0.4574 - accuracy: 0.8594
Epoch 00020: val_accuracy did not improve from 0.89489
62268/62268 [=====] - 258s 4ms/sample - loss: 0.4574 - accuracy: 0.8594

```

```

In [26]: df_hist = pd.DataFrame(history.history)
         plot_results(df_hist)

```



```

In [27]: def get_test_accuracy(model, X_test, y_test):
         test_loss, test_acc = model.evaluate(x=X_test, y=y_test, verbose=0)
         print('accuracy: {acc:0.3f}'.format(acc=test_acc))
         print('loss: {loss:0.3f}'.format(loss=test_loss))

```

```
get_test_accuracy(cnn_model, X_test, y_test)
```

```
accuracy: 0.886
```

```
loss: 0.395
```

## 1.5 4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the label with maximum probability.

```
In [28]: MLP_model = load_model(mlp_checkpoint_best)
        CNN_model = load_model(cnn_checkpoint_best)
```

```
In [29]: num_examples = X_test.shape[0]
        sample = np.random.choice(range(num_examples), size=6)
        X6_test = X_test[sample]
        y6_true_labels = y_test[sample]
        y6_pred_mlp = MLP_model.predict(X6_test)
        y6_pred_cnn = CNN_model.predict(X6_test)
```


```
In [30]: def plot_predictions(n_rows, n_cols, test_images, y_true, y_pred):
```

```
    plt.figure(figsize=(20, 15))
    n = 0
    for i in range(1, n_rows*n_cols+1): # Displaying 6 samples
        if i%2 != 0:
            plt.subplot(n_rows, n_cols, i)
            plt.xticks([])
            plt.yticks([])
            plt.imshow(test_images[n][:,:,0], cmap='gray') # Plot image selected from
            true_label = y_true[n]
            plt.title('True {}'.format(true_label)) # Displaying Label

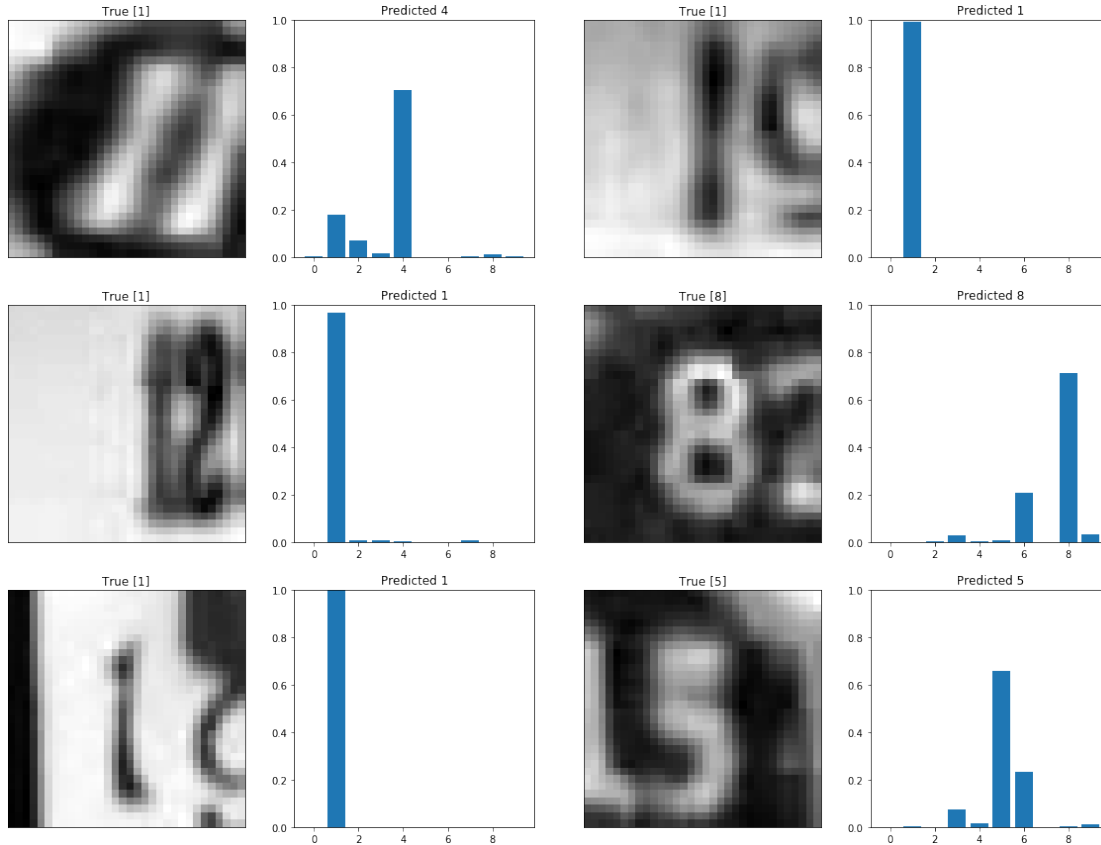
        else:
            plt.subplot(n_rows, n_cols, i)
            plt.bar(range(10), y_pred[n].flatten())
            plt.ylim([0, 1])
            pred_label = np.argmax(y_pred[n])
            plt.title('Predicted {}'.format(pred_label)) # Displaying Label
            n += 1

    plt.show()
```

```
In [31]: plot_predictions(3, 4, X6_test, y6_true_labels, y6_pred_mlp)
```



```
In [32]: plot_predictions(3, 4, X6_test, y6_true_labels, y6_pred_cnn)
```



```
In [33]: y_pred_mlp = MLP_model.predict(X_test)
         y_pred_cnn = CNN_model.predict(X_test)

In [34]: # Convert the true labels and predicted labels into binary matrices,
         # where the hot index (ones) indicates the actual class
         y_pred_mlp_bin = np.zeros(y_pred_mlp.shape)
         y_pred_mlp_bin[np.arange(len(y_pred_mlp)), y_pred_mlp.argmax(1)] = 1

         y_pred_cnn_bin = np.zeros(y_pred_cnn.shape)
         y_pred_cnn_bin[np.arange(len(y_pred_cnn)), y_pred_cnn.argmax(1)] = 1

         mlb = MultiLabelBinarizer()
         y_test_bin = mlb.fit_transform(y_test)

In [35]: def cm_plot(y_true, y_pred, labels, ymap=None, cmap='BuPu', figsize=(16,12)):
         if ymap is not None:
             y_pred = [ymap[yi] for yi in y_pred]
             y_true = [ymap[yi] for yi in y_true]
             labels = [ymap[yi] for yi in labels]
         cm = confusion_matrix(y_true, y_pred, labels=labels)
         cm_sum = np.sum(cm, axis=1, keepdims=True)
```

```

cm_perc = cm / cm_sum.astype(float) * 100
annot = np.empty_like(cm).astype(str)
nrows, ncols = cm.shape
for i in range(nrows):
    for j in range(ncols):
        c = cm[i, j]
        p = cm_perc[i, j]
        if i == j:
            s = cm_sum[i]
            annot[i, j] = '%.1f%%\n%d/%d' % (p, c, s)
        elif c <= 0.5:
            annot[i, j] = ''
        else:
            annot[i, j] = '%.1f%%\n%d' % (p, c)
cm = pd.DataFrame(cm, index=labels, columns=labels)
cm.index.name = 'Real'
cm.columns.name = 'Prediction'
fig, ax = plt.subplots(figsize=figsize)
sns.heatmap(cm, annot=annot, fmt='', ax=ax, linewidth=0.2, cmap = cmap)
plt.ylim(len(cm)+0.05, -0.05)
plt.show()

```

In [36]: `cm_plot(y_test_bin.argmax(axis=1), y_pred_cnn_bin.argmax(axis=1), labels=None, ymap=None, cmap='terrain_r', figsize=(15,9))`

